

Hey everyone! Let's talk about how we built our medical chatbot. It's designed to answer both medical and general questions, and here's how we made it happen under the hood.

At its core, we used Retrieval-Augmented Generation, which combines a Hugging Face model with external knowledge retrieval. To make sure the chatbot has the latest medical info, we web-scraped trusted sources using BeautifulSoup to build a solid knowledge base.

Then, we connected everything with LangChain, so the chatbot can handle complex questions smoothly.

On the backend, we set up an API using Flask, and the frontend uses Axios to communicate with it. This keeps things fast and responsive for users.

Oh, and for transparency purposes, we utilized Claude AI in the debugging and development processes.

So yeah, our chatbot is powered by RAG, Hugging Face, BeautifulSoup, LangChain, Flask, and Axios—all working together to give you smart, reliable answers.